

Progress Report

Week of August 3, 2020

Edris Qarghah



Enrico Fermi Institute
University of Chicago

Weekly Goals

- Finalize tweaking of network distributions
- Create means of visualizing individual path changes
- Observe impacts of breaking an edge:
 - Use traceroutes to identify path changes
 - Identify edges that are no longer being used
 - Infer identity of broken edge
 - Display impact on latency/packetloss
- Observe impact of breaking multiple edges

Daily Log

Monday, August 3

- Network Analytics Discussion meeting.
- Made temporary histograms to confirm distribution matching.
 - Loaded histogram data to dataframes and scaled.
 - Tried alternative means of calculating latency (using the distribution).
 - Adjust packet-loss to account for the probable number of multiplications.
 - Added path hops, latency and packetloss to path lengths function.
 - Updated OWD distribution to account for offset.
- Finished distribution matching (got Ilija's sign-off).
- Fixed issue with lists containing one tuple being loaded as just a tuple.

Tuesday, August 4

- Added function to change edge values (i.e., latency, packetloss).
- Manually changed latency and noted change in shortest paths.
- Added shortest paths dictionary to the graph properties.
 - Encountered issues saving and loading shortest path dictionaries (GML does not like taking non-strings as dictionary keys).
- Created means of converting a shortest path list to a list of edges.
- Created preliminary means of getting changes to paths.
- Added capacity to only draw paths connecting specific PS pairs.

Wednesday, August 5

- Added helper functions to address issue with non-string keys (i.e., pair to string, string to pair).
- Improved graph path change function to account for bi-directionality of edges.
- Met with Ilija about results of killing individual edges. Suggestions:
 - Kill edges entirely (rather than simply increasing latency)
 - Systematically remove each individual edge in a graph (this may break the graph into multiple components).
- Prepared slide for OSG update meeting.
- Read about [time clock sync issues in blockchain and distributed systems](#).
- Ensured that path comparisons still functioned if path could not be completed after edge removal.

Thursday, August 6

- Reviewed OSG update meeting presentation.
- SAND Informal Working Meeting.
 - Learned more about what the others are doing:
 - * Tommy (UMich undergrad) is working on Kibana dashboards for network monitoring
 - * Manjari (UMich undergrad) is working on a UI for network topology analysis (highly relevant to my work)
 - * Petya (U Plovdiv PhD candidate) is working on a Plotly Dash ML web app to provide site summaries.
 - Tommy helped Petya with data pulling issues in Elasticsearch, which was instructive (she was pulling a set of IPs first as a source, then as a destination but was mysteriously getting the same result both ways).
- Wrote functions to systematically remove each edge in a network then tally the resulting path changes.
- Created histograms to display edge ambiguity.
- Modified packetloss histogram to allow comparison of multiple iterations of graphs.
- Updated hops histogram to actually be a histogram!
- Patched issue with histogram obscuring original distribution by setting alpha to .5. A better solution needs to be found.

Friday, August 7

- Finished reports.
- Caught an exception related to how shortest paths were being stored.
- When getting path changes, fixed how edges were being removed from the changed_edges dictionary.
- Updated graph drawing:
 - Adjusted perfsonar node color to make them easier to see (blue instead of brown).
 - Changed the default edge label to None (easier to see paths/colors that way).
- Add naming and autosaving of histograms.
- Updated slides following meeting to clear up ambiguity and structure them better.
- Added details about the process for determining deleted edges:
 1. Determine shortest paths between all PerfSONAR (PS) nodes in G.
 2. Delete an edge and save the network as a new graph, H.
 3. Determine shortest paths between PS nodes in H.
 4. Compare shortest paths between graphs:
 - (a) Record edges in each G path that aren't in the corresponding H path (i.e., edges that are potentially the removed edge).
 - (b) Confirm those edges were removed entirely from H (i.e., they are not on some other path in H).
 - (c) Determine how many different paths each edge was removed from.
 - i. It's likely that the edge removed from the most paths is the deleted edge.
 - ii. If multiple edges were removed an equal number of times, there is ambiguity as to which was deleted.
 5. Repeat 2-4 for every edge in G.
- Caught and corrected issue with edge ambiguity (wasn't fully excluding edges that appear elsewhere in the output graph). This involved a significant rewrite of get_path_changes.
- Fixed font size on histograms.

Achievements

I learned about...

- The “clock problem” for blockchain/“timechain” and distributed systems.
- Other students on the SAND team and how their work relates:
 - Tommy (UMich undergrad) is working on Kibana dashboards for network monitoring
 - Manjari (UMich undergrad) is working on a UI for network topology analysis (highly relevant to my work)
 - Petya (U Plovdiv PhD candidate) is working on a Plotly Dash ML web app to provide site summaries.

I created...

- Means of identifying changed paths:
 - Edges difference between old/new trace-routes.
 - Accounting for lack of edge directionality.
- Means of narrowing down candidate edges:
 - Discarding edges that remain on other paths.
 - Discarding edges not removed from all paths (only works if we know only one. edge)
- Means of visualizing changes:
 - Graphs highlighting only relevant paths.
 - Histograms of candidate edges and impact to distributions.

Roadblocks

Questions

- How can I narrow down the number of potentially impacted edges?
- Is there a way to leverage edge adjacency to help make these determinations?

Problems

- Edges are indexed as tuples, making their use as keys in dictionaries difficult (ordering and saving as strings).
- The toy network has ambiguity with even very limited curated data, which will only be worse with real data.

Challenges

- Highlighting changes to a graph without things getting too cumbersome.
- Developing ways to summarize the impact of changes over many iterations that were intuitive/easy to understand.

Plans for Next Week

- Toy Network
 - Further static network tomography:
 - * Dead vs. throttled edges
 - * Adjusting for change to variable number of edges
 - * Add bandwidth and measure bandwidth impact
 - * Edge removal heatmap
 - Simulate network activity
 - Measure activity from specific nodes/generate Time Series
 - Develop tools for tomography and anomaly detection
- Real Network
 - Take snapshots of traceroutes at various times for comparison.
 - Try to determine whether traceroute variation can be used in the same manner as in toy network to identify dead/ailing edges.